

Results

```
IBM Watson Studio Search in your workspaces Buy Akansha Rawat's Account Frankfurt
```

Projects / Bio Age / Epigenetic_Clock

```
In [4]: # Loading methylation dataset using Pandas
import os, types
import pandas as pd
from boto3.client import Config
import ibm_boto3

def __iter__(self): return 0

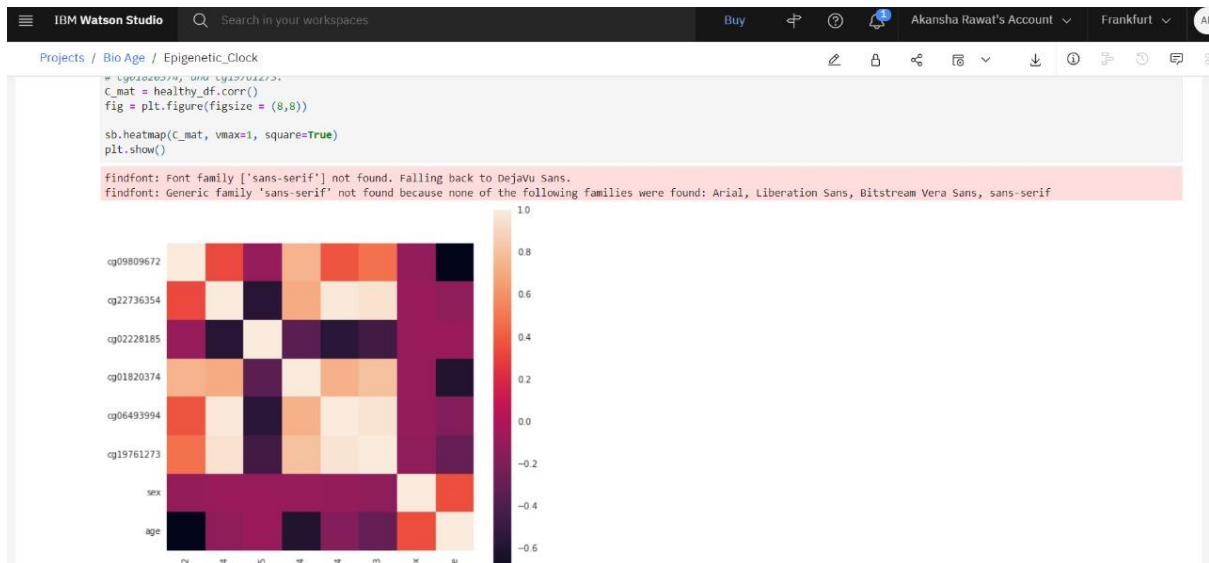
#@hidden_cell
# The following code accesses a file in your IBM Cloud Object Storage. It includes your credentials.
# You might want to remove those credentials before you share the notebook.
client_73f80164cc8d4453ace2ccaeeb9f5d6 = ibm_boto3.client(service_name='s3',
    ibm_api_key_id='VYXxHk8wJtTbga8S5e0x4x5v6WDbYiZ0zWdyeuLSHa8',
    ibm_auth_endpoint='https://iam.cloud.ibm.com/oidc/token',
    config=Config(signature_version='oauth'),
    endpoint_url='https://s3.private.eu.cloud-object-storage.appdomain.cloud')

body = client_73f80164cc8d4453ace2ccaeeb9f5d6.get_object(Bucket='bioage-donotdelete-pr-re6stby4ah9ab4',Key='Healthy_Methylation_Dataset.csv')['Body']
# add missing __iter__ method, so pandas accepts body as file-like object
if not hasattr(body, "__iter__"): body.__iter__ = types.MethodType(__iter__, body)

healthy_df = pd.read_csv(body)
healthy_df.head()
```

Out[4]:

		...1	cg09809672	cg22736354	cg02228185	cg01820374	cg06493994	cg19761273	sex	age
0	GSM507152		0.392464	0.281630	0.352591	0.315924	0.111604	0.211683	1	71.06
1	GSM507153		0.377919	0.238900	0.520396	0.323641	0.127965	0.184307	1	69.65
2	GSM507154		0.458522	0.204078	0.519273	0.363714	0.119891	0.215067	1	65.79



```
[8]: # Normalizing the methylation and sex data with a Standard Scaler.
X = healthy_df[['cg09809672', 'cg22736354', 'cg02228185', 'cg01820374', 'cg06493994', 'cg19761273', 'sex']]

# Separating X vs. y dataframes
X_std = pd.DataFrame(std_scaler.fit_transform(X), columns=X.columns)
y = healthy_df['age']

[9]: # Separating dataset into train and test subsets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.33, random_state = 42)
```

```
In [5]: # Shuffle dataframe to randomize data order, possibly preventing confounding factors
healthy_df = shuffle(healthy_df)

# Remove patient ID column
healthy_df = healthy_df.drop(['...1'], axis=1)

# Drop all rows with NaN values
healthy_df = healthy_df.dropna()

# Reset Index
healthy_df.reset_index(inplace=True, drop=True)

healthy_df.head()
```

```
Out[5]:
```

	cg09809672	cg22736354	cg02228185	cg01820374	cg06493994	cg19761273	sex	age
0	0.861134	0.030339	0.772940	0.764273	0.038618	0.407311	1	0.00
1	0.665470	0.111330	0.644090	0.528570	0.050360	0.325280	0	4.25
2	0.613737	0.231627	0.798076	0.392510	0.179445	0.307907	0	32.00
3	0.847000	0.076200	0.781000	0.762000	0.030100	0.416000	0	0.00
4	0.875370	0.127060	0.704220	0.589560	0.014410	0.668570	1	0.00

```
In [6]: # Checking if there are any remaining NaNs in the dataset
np.where(pd.isnull(healthy_df))
```

```
Out[6]: (array([], dtype=int64), array([], dtype=int64))
```

```
: predictions = random_forest_regressor.predict(X_test)

# Since age cannot be negative, changing all negative predictions to age 0
for n, element in enumerate(predictions):
    if element < 0:
        predictions[n] = 0

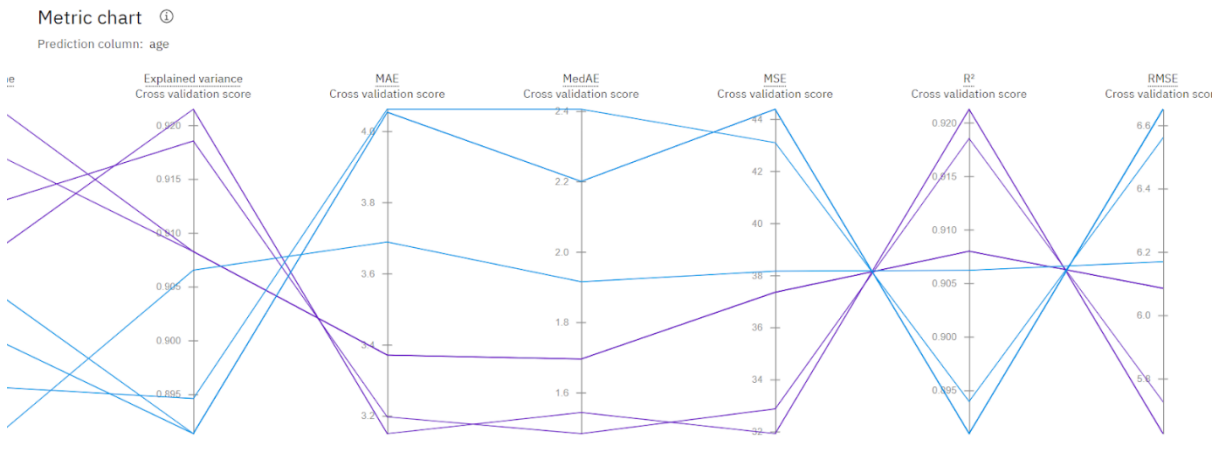

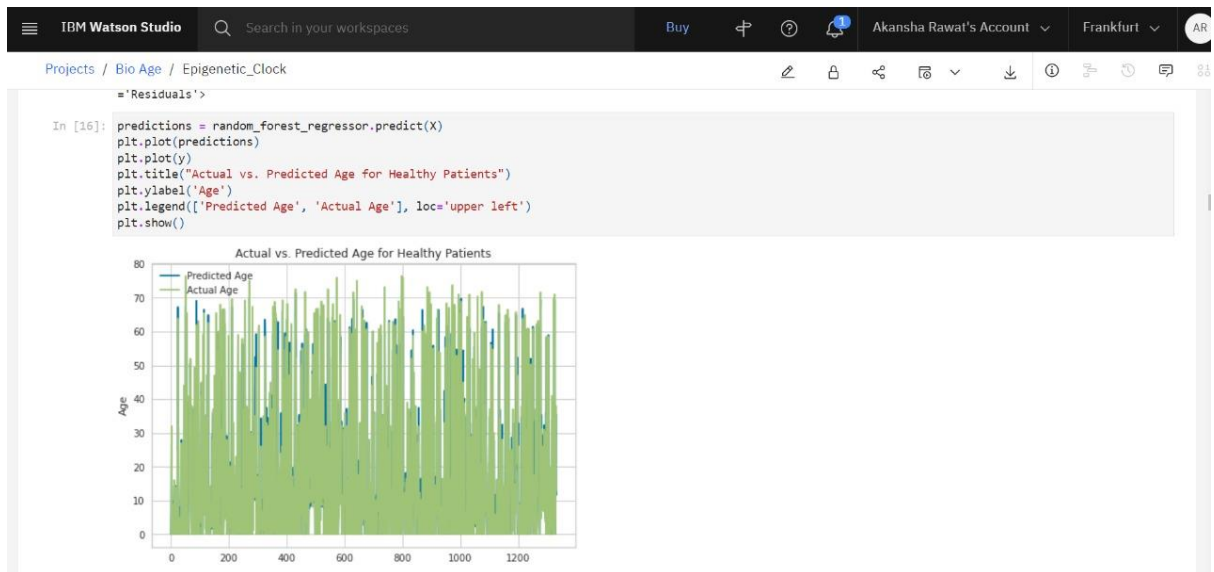
# Looking at sample predictions for the testing set
for i in range(0, 10):
    print("Prediction:", predictions[i], "\tActual:", y_test.iloc[i])
```

```
Prediction: 67.63982352941176 Actual: 68.2
Prediction: 0.754901960764706 Actual: 0.0
Prediction: 61.47823529411764 Actual: 60.9
Prediction: 17.63235294117647 Actual: 17.0
Prediction: 0.0 Actual: 0.0
Prediction: 8.235294117352941 Actual: 12.91666667
Prediction: 13.014705882764705 Actual: 13.41666667
Prediction: 30.235294117647058 Actual: 30.0
Prediction: 7.563725491529412 Actual: 5.91666667
Prediction: 0.0 Actual: 0.0
```

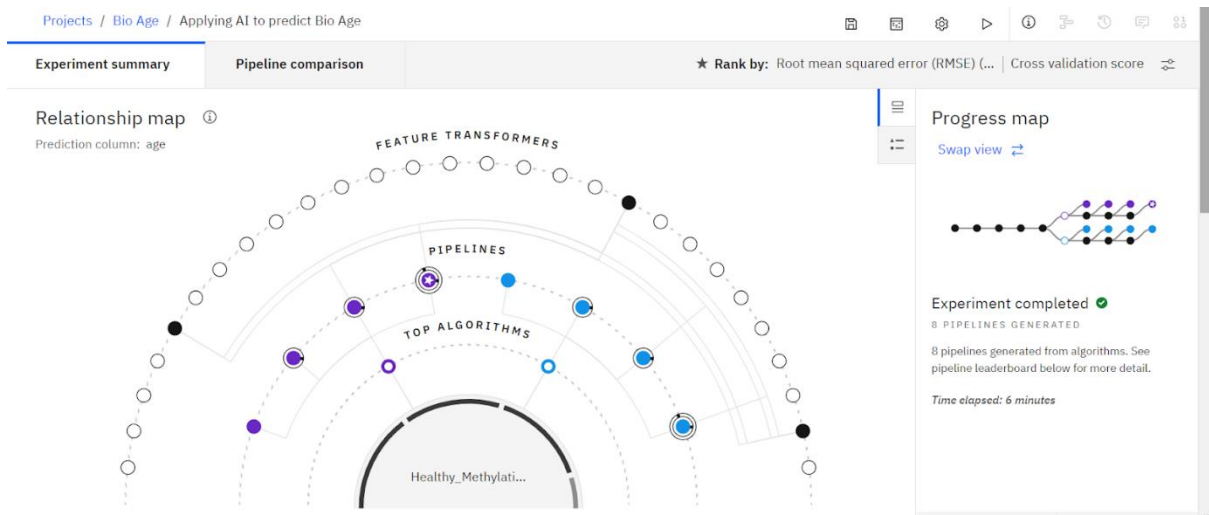
```
# Building and training the Random Forest Regressor model
# Optimal value for n_estimators was determined by trial and error, comparing the score for each trial
random_forest_regressor = RandomForestRegressor(n_estimators = 17, random_state = 0)
random_forest_regressor.fit(X_train, y_train)

# Accuracy on the testing set
test_acc = random_forest_regressor.score(X_test, y_test)
print(test_acc)
```

```
0.9302161785526581
```



	Rank ↑	Name	Algorithm	Specialization	RMSE (Optimized) Cross Validation	Enhancements	Build time
★	1	Pipeline 4	Random Forest Regressor		5.626	HPO-1 FE HPO-2	00:02:06
	2	Pipeline 3	Random Forest Regressor		5.727	HPO-1 FE	00:00:46
	3	Pipeline 2	Random Forest Regressor		6.085	HPO-1	00:00:18
	4	Pipeline 1	Random Forest Regressor		6.085	None	00:00:01
	5	Pipeline 8	LGBM Regressor		6.169	HPO-1 FE HPO-2	00:02:10
	6	Pipeline 7	LGBM Regressor		6.560	HPO-1 FE	00:01:00



Promote to space

Use a deployment space to organize supporting resources such as input data and environments; deploy models or functions to generate predictions or solutions; and view or edit deployment details.

Target space

Models

Tags (optional)

Start typing to add tags

☐ Go to the model in the space after promoting it

Selected assets (1)

Asset name	Format
Applying AI to predict Bio Age - P4 Random Forest Regressor	Model

Select version

Promoting a version of an asset to a space creates a new asset

Cancel Promote

Create a deployment

Associated asset

Applying AI to predict Bio Age - P4 Random Forest Regressor

Deployment type

Online
Run the model on data in real-time, as data is received by a web service.

Batch
Run the model against data as a batch process.

Name

Bio Age Pred

Cancel Create

IBM Watson Studio

Search in your workspaces

Buy

Akansha Rawat's Account

Frankfurt

AR

Deployments / Models / Applying AI to predict Bio Age - P... /

Bio Age Pred

Deployed

Online

API reference

Test

Enter input data

Input

Paste JSON

Enter data manually or use a CSV file to populate the spreadsheet. Max file size is 50 MB.

Download CSV template
Browse local files
Search in space
Clear all

	...1 (other)	cg09809672 (double)	cg22736354 (double)	cg02228185 (double)	cg01820374 (double)	cg06493994 (double)	cg19761273 (double)
2	GSM50715	0.37791864	0.238899924	0.520396483	0.323641214	0.127964701	0.184306617

1,440 rows, 8 columns

Predict

Browse local files to add your dataset file and click predict.

Prediction results

Regression classification

Prediction distribution

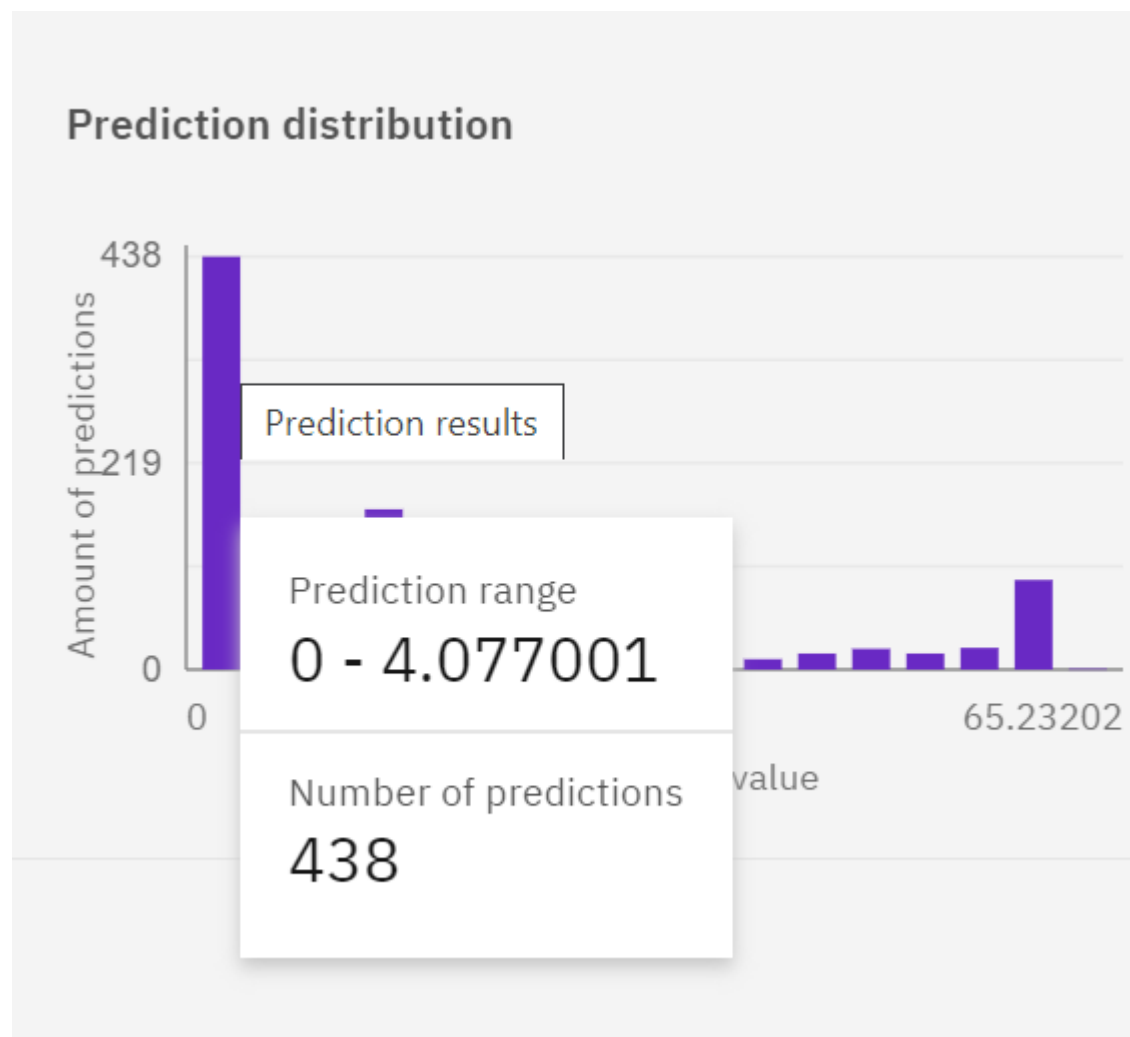
	Prediction
1	65.0582871954125
2	64.63239600859492
3	63.4382524835058
4	52.52556903678251
5	62.796697088034755
6	64.83322695077183
7	58.349404116710986

Download

Bio Age Pred_test...json

Show all

After predicting, download the results as a .json file.



Each bar shows the range and estimated predictions.